

## Keith Winstein — Ph.D. Thesis Proposal

**Thesis title:** Transport Architectures for an Evolving Internet

In the Internet architecture, transport protocols are the glue between an application’s needs and the network’s abilities. The Transmission Control Protocol, or TCP, plays this role for the vast majority of Internet traffic, and TCP’s implementation may be the most widely used computer program in the world. Viewed through a broader lens, the Hypertext Transfer Protocol (HTTP) over TCP is the true transport for most applications—including the World Wide Web, smartphone and tablet apps, and video services like YouTube, Netflix, and Hulu.

As the Internet has matured over the last 30 years into a global utility, HTTP and TCP have proven extraordinarily successful. But as the Internet’s underlying networks have evolved, the transport layer has begun showing its age. The implicit assumptions of these protocols—and their broad applicability to most applications—have held less and less well.

TCP was designed for a world where users have one endpoint address and continual connectivity to a wired network with “dumb” gateways with limited memory, a network whose only source of packet loss is from gateways that run out of memory to store packets in flight and whose only variation in performance comes from cross traffic. In TCP’s world, applications have one objective: the speedy completion of a long-running reliable file transfer.

Not one of those design assumptions still holds. The performance of wireless networks varies greatly in time, even in the absence of other users. Packet loss may occur for reasons other than gateway buffer overflow. Users regularly “roam” among networks (e.g. from one Wi-Fi network to another, or to a cellular network), changing their IP addresses in the process. Laptops and smaller computing devices go to sleep or lose Internet access for hours at a time. Gateways and routers have considerable sophistication, large buffer memories, and complex behavior.

Furthermore, long-running file transfers now constitute a minority of application demands on the network. Today’s network applications represent a diverse menagerie of desires. Remote procedure calls and user interface tools (e.g. SSH, VNC) may send only a trickle of traffic, but care greatly about latency. Web pages generally care about the completion time of an ensemble of small file transfers, often sent from independent servers. Videoconferencing programs would like a compromise between high throughput and low delay. Batch processing applications like MapReduce may care about the tail of the distribution of completion time of a transaction. Applications may ask TCP to send information that becomes superseded or obsolete before it is successfully acknowledged—meaning TCP should stop trying to deliver the data reliably.

The widespread use of TCP has meant that network technologies now must be designed to satisfy TCP’s implicit assumptions about them. With my advisor, Hari Balakrishnan, I earlier argued in a position paper<sup>1</sup> that this constraint has hampered the evolution of the Internet, which would be freer to evolve if the transport layer’s assumptions were made explicit, and the protocol made a function of those assumptions.

Working with collaborators at MIT, I have built three systems that explore this type of design:

1. **Remy**, a tool that generates end-to-end congestion-control schemes automatically as a function of a designer’s goals and assumptions about the network, allowing the transport layer to evolve with the network. By asking Remy to search for the best schemes subject to different constraints, the tool helps investigate the fundamental limits of decentralized control on a packet-switched network.
2. **Sprout**, a protocol for interactive real-time applications such as a videoconference, that achieves many-fold gains in throughput and delay compared with Skype, Facetime, Google Hangouts, and TCP congestion-control schemes
3. **Mosh/SSP**, a remote terminal application and protocol that supports roaming, intermittent connectivity, and rolling latency compensation

---

<sup>1</sup>K. Winstein and H. Balakrishnan, End-to-End Transmission Control by Modeling Uncertainty about the Network State, in Proceedings of ACM HotNets 2011, Cambridge, Mass., November 2011




Figure 1: Remy designs congestion-control schemes automatically to achieve desired outcomes. The algorithms it produces may replace the congestion-control module of a TCP implementation, and fit into a network library or kernel module that implements congestion control (DCCP, SCTP, the congestion manager, application-layer transmission control libraries, ns-2 modules, etc.).

Based on this experience, my thesis is that the explicit modeling of assumptions about network behavior and user objectives by the transport layer allows the creation of new application behaviors, frees network technologies to evolve, and improves the experience of real-world applications.

## 1. Remy: computer-generated congestion control

*(This section will be based on material published in K. Winstein and H. Balakrishnan, TCP ex Machina: Computer-Generated Congestion Control, in Proc. ACM SIGCOMM 2013, Hong Kong, China, August 2013.)*

Is it possible for a computer to discover the right rules for congestion control in heterogeneous and dynamic networks? Should computers, rather than humans, be tasked with developing congestion control methods? And just how well can we make computers perform this task?

We investigated these questions and found that computers can design schemes that in some cases surpass the best human-designed methods to date, when supplied with the appropriate criteria by which to judge a congestion-control algorithm. In my dissertation, I will discuss how this style of transport-layer protocol design can give more freedom to network architects and link-layer designer, and provide angles for understanding fundamental questions about the limits of decentralized network algorithms.

Congestion control, a fundamental problem in multi-user computer networks, addresses the question: when should an endpoint transmit each packet of data? An ideal scheme would transmit a packet whenever capacity to carry the packet was available, but because there are many concurrent senders and the network experiences variable delays, this question isn't an easy one to answer. On the Internet, the past thirty years have seen a number of innovative and influential answers to this question, with solutions embedded at the endpoints (mainly in TCP) aided occasionally by queue management and scheduling algorithms in bottleneck routers that provide signals to the endpoints.

This area has continued to draw research and engineering effort because new link technologies and subnetworks have proliferated and evolved. For example, the past few years have seen an increase in wireless networks with variable bottleneck rates; datacenter networks with high rates, short delays, and correlations in offered load; paths with excessive buffering (now called “bufferbloat”); cellular wireless networks with highly variable, self-inflicted packet delays; links with non-congestive stochastic loss; and networks with large bandwidth-delay products. In these conditions, the classical congestion-control methods embedded in TCP can perform poorly, as many papers have shown.

Without the ability to adapt its congestion-control algorithms to new scenarios, TCP's inflexibility con-

strains architectural evolution, as we noted in the 2011 position paper.<sup>2</sup> Subnetworks and link layers are typically evaluated based on how well TCP performs over them. This scorecard can lead to perverse behavior, because TCP’s network model is limited. For example, because TCP assumes that packet losses are due to congestion and reduces its transmission rate in response, some subnetwork designers have worked hard to hide losses. This often simply adds intolerably long packet delays. One may argue that such designs are misguided, but the difficulties presented by “too-reliable” link layers have been a perennial challenge for 25 years and show no signs of abating. With the rise of widespread cellular connectivity, these behaviors are increasingly common and deeply embedded in deployed infrastructure.

The designers of a new subnetwork may well ask what they should do to make TCP perform well. This question is surprisingly hard to answer, because the so-called teleology of TCP is unknown: exactly what objective does TCP congestion control optimize? TCP’s dynamic behavior, when competing flows enter and leave the network, remains challenging to explain. In practice, the need to “make TCP perform well” is given as a number of loose guidelines, such as IETF RFC 3819, which contains dozens of pages of qualitative best current practice. The challenging and subtle nature of this area means that the potential of new subnetworks and network architectures is often not realized.

## Design overview

How should we design network protocols that free subnetworks and links to evolve freely, ensuring that the endpoints will adapt properly *no matter what* the lower layers do? We believe that the best way to approach this question is to take the design of specific algorithmic mechanisms out of the hands of human designers (no matter how sophisticated!), and make the end-to-end algorithm be a function of the desired overall behavior.

As with Mosh and Sprout, we start by explicitly stating an objective for congestion control; for example, given an unknown number of users, we may optimize some function of the per-user throughput and packet delay, or a summary statistic such as average flow completion time. Then, instead of writing down rules by hand for the endpoints to follow, we start from the desired objective and work backwards in three steps:

1. First, model the protocol’s prior assumptions about the network; i.e., the “design range” of operation. This model may be different, and have different amounts of uncertainty, for a protocol that will be used exclusively within a data center, compared with one intended to be used over a wireless link or one for the broader Internet. A typical model specifies upper and lower limits on the bottleneck link speeds, non-queueing delays, queue sizes, and degrees of multiplexing.
2. Second, define a traffic model for the offered load given to endpoints. This may characterize typical Web traffic, video conferencing, batch processing, or some mixture of these. It may be synthetic or based on empirical measurements.
3. Third, use the modeled network scenarios and traffic to design a congestion-control algorithm that can later be executed on endpoints.

We have developed an optimization tool called Remy that takes these models as input, and designs a congestion-control algorithm that tries to maximize the total expected value of the objective function, measured over the set of network and traffic models. The resulting pre-calculated, optimized algorithm is then run on actual endpoints; no further learning happens after the offline optimization. The optimized algorithm is run as part of an existing TCP sender implementation, or within any congestion-control module. No receiver changes are necessary (as of now).

## Summary of results

We have implemented Remy. Running on a 48-core server at MIT, Remy generally takes a few hours of wall-clock time (one or two CPU-weeks) to generate congestion-control algorithms offline that work on a wide range of network conditions.

---

<sup>2</sup>Above, footnote 1.

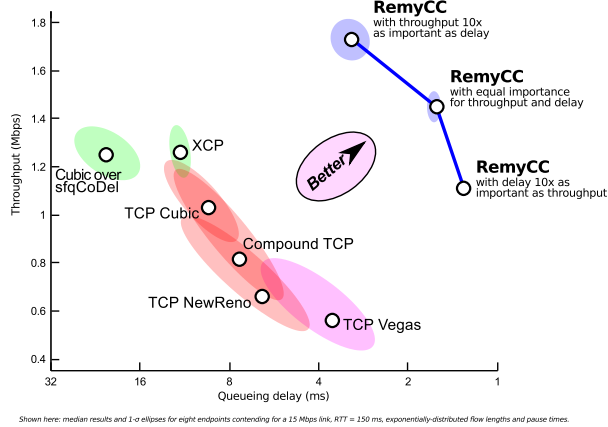


Figure 2: Results for each of the schemes over a 15 Mbps dumbbell topology with  $n=8$  senders, each alternating between flows of exponentially-distributed byte length (mean 100 kilobytes) and exponentially-distributed off time (mean 0.5 s). Medians and  $1-\sigma$  ellipses are shown. The blue line represents the efficient frontier, which here is defined entirely by the RemyCCs.

Our main results from several simulation experiments with Remy are as follows:

1. For networks broadly consistent with the assumptions provided to Remy at design time, the machine-generated algorithms dramatically outperform existing methods, including TCP Cubic, Compound TCP, and TCP Vegas.
2. Comparing Remy’s algorithms with schemes that require modifications to network gateways, including Cubic-over-sfqCoDel and XCP, Remy generally matched or surpassed these schemes, despite being entirely end-to-end.
3. We measured the tradeoffs that come from specificity in the assumptions supplied to Remy at design time. As expected, more-specific prior information turned out to be helpful when it was correct, but harmful when wrong. We found that RemyCC schemes performed well even when designed for an order-of-magnitude variation in the values of the underlying network parameters.

On a simulated 15 Mbps fixed-rate link with eight senders contending and an RTT of 150 ms, a computer-generated congestion-control algorithm achieved the following improvements in median throughput and reductions in median queueing delay over these existing protocols. The results are plotted in Figure 2.

Protocol	Median speedup	Median delay reduction
Compound	2.1×	2.7×
NewReno	2.6×	2.2×
Cubic	1.7×	3.4×
Vegas	3.1×	1.2×
Cubic/sfqCoDel	1.4×	7.8×
XCP	1.4×	4.3×

In a trace-driven simulation of the Verizon LTE downlink with four senders contending, the *same* computer-generated protocol achieved these speedups and reductions in median queueing delay:

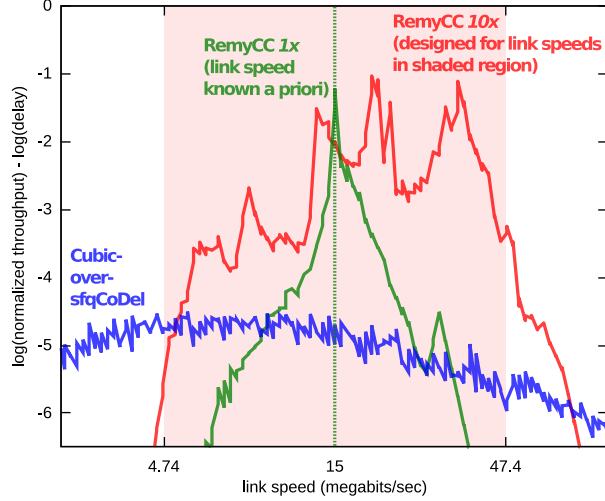


Figure 3: How helpful is prior information about the network? Here we show the performance of two end-to-end RemyCCs that were designed with different prior information about the network, compared with an in-network algorithm (Cubic-over-sfqCoDel), as the link speed varies. Despite running only at the sender, the RemyCCs each outperform Cubic-over-sfqCoDel over almost their entire design ranges. But when a RemyCC’s assumptions aren’t met, performance deteriorates. The extent to which there is a tradeoff between the “operating range” and performance of a RemyCC is so far unresolved.

Protocol	Median speedup	Median delay reduction
Compound	1.3×	1.3×
NewReno	1.5×	1.2×
Cubic	1.2×	1.7×
Vegas	2.2×	0.44× ↓
Cubic/sfqCoDel	1.3×	1.3×
XCP	1.7×	0.78× ↓

## Remy as a tool for network science

In addition to its utility as a tool for constructing congestion-control algorithms, Remy is an aid for understanding the limits and structure of distributed networking problems. In my dissertation, I will report the results of our ongoing computational experiments that use Remy to address several open questions in network science and the optimization of decentralized algorithms:

- How helpful is prior information about a network? To what extent are the RemyCCs’ gains (relative to human-designed algorithms) due to the benefits of explicit optimization towards an objective, vs. simply to the narrowed “operating range” implied by the stated assumptions? Is there a tradeoff between operating range and performance—in other words, if we design a RemyCC for a broader and broader range of network conditions, does its performance necessarily get worse and worse (Figure 3)? Along the same lines: *can* we successfully design a RemyCC that still beats TCP over a 10,000-fold range of throughputs or round-trip times?
- To what extent is the end-to-end congestion-control problem decomposable? If we wish to design a RemyCC that works well on a complex network (e.g. the actual Internet), must we also optimize on a model of that complex network? What is the penalty associated with optimizing for a simple model, but then running a computer-engineered solution in the real world? How can we explain or put bounds on this cost, of using a simple model in a complex world?

- Which congestion signals (in machine-learning terms, which features) are helpful in different network scenarios (e.g. high multiplexing, high BDP) and which aren't? Can we develop a general theory to predict which features will be helpful, when?
- What is the “cost of compatibility”? How much does it hurt the performance of a RemyCC if we demand that it also achieve the right result when competing against “buffer-filling” TCP algorithms?
- How can we quantify the “cost of coexistence”? How well can we co-optimize a pair of RemyCCs that satisfy different objectives, so as to achieve an appropriate division of network resources—even when each RemyCC is uncertain about how many other RemyCCs are contending for the same bottleneck, or what objective they are each trying to optimize.

## 2. Sprout: using forecasts of network variation to achieve an explicit objective

*(This section will be based on material published in K. Winstein, A. Sivaraman, and H. Balakrishnan, Stochastic Forecasts Achieve High Throughput and Low Delay over Cellular Networks, in Proc. USENIX NSDI 2013, Lombard, Ill., April 2013.)*

Cellular wireless networks have become a dominant mode of Internet access. These mobile networks, which include LTE and 3G (UMTS and 1xEV-DO) services, present new challenges for network applications, because they behave differently from wireless LANs and from the Internet’s traditional wired infrastructure.

Cellular wireless networks experience rapidly varying link rates and occasional multi-second outages in one or both directions, especially when the user is mobile. As a result, the time it takes to deliver a network-layer packet may vary significantly, and may include the effects of link-layer retransmissions. Moreover, these networks schedule transmissions after taking channel quality into account, and prefer to have packets waiting to be sent whenever a link is scheduled. They often achieve that goal by maintaining deep packet queues. The effect at the transport layer is that a stream of packets experiences widely varying packet delivery rates, as well as variable, sometimes multi-second, packet delays.

For an interactive application such as a videoconferencing program that requires both high throughput and low delay, these conditions are challenging. If the application sends at too low a rate, it will waste the opportunity for higher-quality service when the link is doing well. But when the application sends too aggressively, it accumulates a queue of packets inside the network waiting to be transmitted across the cellular link, delaying subsequent packets. Such a queue can take several seconds to drain, destroying interactivity (Figure 4).

Experiments with Microsoft’s Skype, Google’s Hangout, and Apple’s Facetime running over traces from commercial 3G and LTE networks show the shortcomings of the transport protocols in use and the lack of adaptation required for a good user experience. The transport protocols deal with rate variations in a reactive manner: they attempt to send at a particular rate, and if all goes well, they increase the rate and try again. They are slow to decrease their transmission rate when the link has deteriorated, and as a result they often create a large backlog of queued packets in the network. When that happens, only after several seconds and a user-visible outage do they switch to a lower rate.

By contrast, *Sprout* is a transport protocol designed to satisfy a particular objective on behalf of interactive applications on variable-quality networks. *Sprout* uses the receiver’s observed packet arrival times as the primary signal to determine how the network path is doing, rather than the packet loss, round-trip time, or one-way delay. Moreover, instead of the traditional reactive approach where the sender’s window or rate increases or decreases in response to a congestion signal, the *Sprout* receiver makes a short-term forecast (at times in the near future) of the bottleneck link rate using probabilistic inference. From this model, the receiver predicts how many bytes are likely to cross the link within several intervals in the near future with at least 95% probability. The sender uses this forecast to transmit its data, bounding the risk that the queuing delay will exceed some threshold, and maximizing the achieved throughput within that constraint.

We conducted a trace-driven experimental evaluation using data collected from four different commercial cellular networks (Verizon’s LTE and 3G 1xEV-DO, AT&T’s LTE, and T-Mobile’s 3G UMTS). We compared




Figure 4: Skype and Sprout on the Verizon LTE downlink trace. For Skype, overshoots in throughput lead to large standing queues. Sprout has an explicit objective: send as much as possible, but keep each packet’s delay less than 100 ms with 95% probability.

Sprout with Skype, Hangout, Facetime, and several TCP congestion-control algorithms, running in both directions (uplink and downlink).

The following table summarizes the average relative throughput improvement and reduction in self-inflicted queueing delay for Sprout compared with the various other schemes, averaged over all four cellular networks in both directions. Metrics where Sprout did not outperform the existing algorithm are highlighted in red:

App/protocol	Avg. speedup vs. Sprout	Delay reduction (from avg. delay)
Sprout	1.0×	1.0× (0.32 s)
Skype	2.2×	7.9× (2.52 s)
Hangout	4.4×	7.2× (2.28 s)
Facetime	1.9×	8.7× (2.75 s)
Compound	1.3×	4.8× (1.53 s)
TCP Vegas	1.1×	2.1× (0.67 s)
LEDBAT	1.0×	2.8× (0.89 s)
Cubic	0.91×	79× (25 s)
Cubic-CoDel <sup>3</sup>	0.70×	1.6× (0.50 s)

The Sprout experience suggests that modeling a network’s state explicitly, albeit with a simplified, imperfect model, and with the aim of achieving a particular objective, can produce many-fold improvements on conventional metrics over contemporary systems.

### 3. Mosh (mobile shell) and the State Synchronization Protocol

(This section will be based on material published in K. Winstein and H. Balakrishnan, *Mosh: An Interactive Remote Shell for Mobile Clients*, in *Proc. USENIX Annual Technical Conference 2012, Boston, Mass., June 2012*.)

Remote terminal applications are almost as old as packet-switched data networks. Starting with RFC 15

in 1969, protocols like TELNET, SUPDUP, and BSD’s rlogin and rsh have played an important role in the Internet’s development. The most popular such application today is the Secure Shell, or SSH.

SSH has two weaknesses that can make it unpleasant. First, because SSH and previous remote terminals run over TCP, they don’t support roaming between IP addresses, or generally preserve sessions when connectivity is intermittent. A laptop will happily suspend for a commute to work, but after wakeup its SSH sessions will have frozen or died.

Second, SSH operates strictly in character-at-a-time mode, with all echoes and line editing performed by the remote host. As a result, its interactive performance can be poor over wide-area wireless (e.g., EV-DO, UMTS, LTE) or transcontinental networks (e.g., to cloud computing facilities or remote data centers), and sessions are almost unusable over paths with non-trivial packet loss.

When loaded or when the signal-to-noise ratio is low, delays on many wireless networks reach several *seconds* because of deep packet queues (“bufferbloat”) or over-zealous link-layer retransmissions. Many home networks also suffer from multi-second delays under load. Trying to type, or correct a typo, over such networks is unpleasant.

Mosh (mobile shell) is addressed at both problems. Mosh is a remote terminal application that supports IP roaming, intermittent connectivity, and marginal network connections. Mosh performs predictive client-side echoing and line editing without any change to server software, and without regard to which application is running. Mosh makes remote servers feel more like the local computer, because most keystrokes are reflected immediately on the user’s display—even in full-screen programs like a text editor or mail reader.

These features are possible because Mosh operates at a different layer from SSH. While SSH conveys an octet-stream to TCP, and eventually to a separate client-side terminal emulator to be interpreted and rendered in cells on the screen, Mosh contains a server-side terminal emulator and uses a new transport protocol—the State Synchronization Protocol—to synchronize terminal screen states over the network, using the principle of application-layer framing.

Because both the server and client maintain an image of the screen state, Mosh can support intermittent connectivity and local editing, and can adjust its network traffic (even superseding screen states that are no longer important) to avoid filling network buffers on slow links. As a result, unlike in SSH, in Mosh “Control-C” always works to cease output from a runaway process within an RTT.

Giving SSP the flexibility to send the packet that best advances the user’s objective—in this case, low UI latency—leads to better application behavior:

- SSP implements “single-packet roaming,” which allows the client to switch public IP addresses and keep its connection by sending one packet from the new address—without needing to know that it has switched IP addresses at all.
- The protocol also uses a novel technique, called “pretransmissions,” to code opportunistically against packet loss. When formulating a transmission to the receiver, the sender will normally use the most recent state transmitted as a predicate (even if this state has not yet been acknowledged, as long as it has not timed out). This is what TCP does, by sending only new data in each segment, and this policy represents a gamble—if the older segment does not arrive, neither transmission will be useful to the receiver without a retransmission.

Instead, SSP may optionally choose to base its transmission only on a state that has actually been acknowledged by the receiver. This will typically require more bits on the wire, but usually by only a modest amount compared with framing overhead. The benefit is that the new transmission is sufficient to update the receiver to the current state by itself, even if previous in-flight transmissions are lost.

- In addition, Mosh introduces the technique of “rolling latency compensation” to improve the user’s experience when interacting with a server-side application across the network. The Mosh client runs a predictive model of the application in the background, and uses the model to do intelligent client-side echoing and line editing. When confident in a prediction, Mosh will display it to the user, but underlined. As predictions are confirmed, Mosh removes the underline. The effect is a sliding “prediction window” of keystrokes shown locally but not yet confirmed by the server.






Figure 5: In 40 hours of user-submitted sessions, Mosh successfully predicted and displayed 70% of user keystrokes immediately. The error rate was less than 1%.

According to data we have collected from contributing users, more than two-thirds of the keystrokes in a typical Unix session can be displayed instantly with a conservative model of application behavior (Figure 5). Mosh’s empirical approach to local echo works in full-screen programs like a text editor or mail reader as well as at the command line.

We have implemented Mosh in C++ and published it as free software. It is in wide deployment and has been packaged for or included in Debian GNU/Linux, Ubuntu, Fedora, Red Hat EPEL, Gentoo, Arch Linux, FreeBSD, NetBSD, OpenBSD, Android, Cygwin, Homebrew, MacPorts, and OpenSolaris.

## 4. Conclusion and Timeline

My thesis is that the explicit modeling of assumptions about network behavior and user objectives by the transport layer allows the creation of new application behaviors, frees network technologies to evolve, and improves the experience of real-world applications.

I propose to defend and deliver the dissertation in time for graduation in June 2014. The Mosh and Sprout sections will largely be based on the published conference papers about those systems. The Remy section will also be based on its conference paper (SIGCOMM 2013), with the addition of our further work on Remy that, along with my co-authors, we are submitting this coming January for review at SIGCOMM 2014.

In addition to the above, I am working on a project to produce an Internet stored-video system (similar to YouTube, Netflix, and Amazon Streaming) from first principles, based again on the principle of explicit objectives and network modeling, as well as a new design for “networked video coding.” If this work is successful, I would plan to include it as a chapter in the dissertation.